

ANNOUNCEMENTS

1

Illinois Passes First State Law Requiring Independent Audits of Frontier AI Labs

Source: *Illinois House, via Transparency Coalition, Axios* · 2026-05-30

The Illinois House passed the AI Safety Measures Act by a vote of 110 to 0. If signed, it would be the first US law to mandate annual third-party audits of frontier AI labs, requiring independent safety audits, incident reporting,...

2

Anthropic Moves Claude Agents Inside the Customer Perimeter With Self-Hosted Sandboxes and Private MCP Tunnels

Source: *Anthropic, Code with Claude London* · 2026-05-19

At its London developer conference, Anthropic released self-hosted sandboxes in public beta and MCP tunnels in research preview. Self-hosted sandboxes let the tool-execution layer of an agent run on infrastructure controlled by the customer.

3

FTC Fines Three Firms \$930,000 for Falsely Marketing an AI Service That Did Not Exist

Source: *Federal Trade Commission* · 2026-05-21

The FTC announced settlements totaling \$930,000 with Cox Media Group, MindSift, and 1010 Digital Works over an advertising product marketed as an AI-powered Active Listening service.

4

DeepSeek Makes Its 75 Percent Price Cut Permanent, Escalating the AI Price War

Source: *Reuters* · 2026-05-23

Chinese AI developer DeepSeek announced it will permanently keep a 75 percent price cut on its flagship V4-Pro model. The change drops V4-Pro pricing to roughly a quarter of its original level, with output token costs cited at around \$0.

WHAT IT MEANS

For a regional bank operating across a multi-state footprint in the northeast and mid-Atlantic, Illinois has just written the template that other state legislatures will copy: mandated outside audits of...

Like:

It is like the first state to require restaurants to post health-inspection grades in the window.

This directly addresses the single most common reason regulated banks have stalled on AI agents: that running tools on a vendor's infrastructure crosses a security boundary that compliance, legal, and...

Like:

It is the difference between letting a contractor take your files back to their office and letting them sit in...

Regulators are now treating false or exaggerated AI capability claims as straightforward deceptive marketing — and the FTC is enforcing.

Like:

It is like selling a home security system with cameras that were never plugged in.

The cost of frontier-grade AI is falling fast and is now being treated as a permanent floor rather than a promotion.

Like:

It is like a new airline permanently pricing every seat at a quarter of the going rate.

WHAT IT MEANS FOR: M&T BANK

REGULATORY COMPLIANCE

M&T Bank's AI vendor stack — from call-center intelligence to AML transaction monitoring to lending-decision support — will soon face a state-level audit mandate that looks a lot like what SOC 2 did for cloud...

TECHNOLOGY ARCHITECTURE

M&T Bank has likely evaluated and shelved multiple AI agent use cases — automated regulatory filing prep, internal policy search for front-line staff, loan-document review — because the vendor-infrastructure security question was a hard no.

VENDOR GOVERNANCE

The FTC action establishes that AI-washing carries real enforcement consequences, and the liability flows to downstream users who repeat vendor claims.

VENDOR ECONOMICS

M&T Bank's AI procurement budget has likely been built on model-pricing assumptions that are now out of date. A permanent 75 percent price cut from a credible frontier competitor resets the entire market.

THE WEEK IN ONE LINE

AI compliance frameworks arrive at the state level, the security barrier that kept agents out of regulated banks gets a direct architectural fix, and the cost curve for AI procurement bends sharply downward just as workforce operating models begin to shift beneath the desktop.