HEADS DOWN TO HEADS UP

Weekly Intelligence Brief

Happenings in AI last week, and what they mean for your business.

ISSUE
Week 25 · 2026

PREPARED FOR	WEEK OF
M&T Bank Regional Bank	**June 15 to June 19, 2026** Prepared by Keel3, LLC

ANNOUNCEMENTS	WHAT IT MEANS	WHAT IT MEANS FOR: M&T BANK
1 **Illinois Passes First State Law Requiring Independent Audits of Frontier AI Labs** *Illinois House, via Transparency Coalition, Axios, May 30, 2026* The Illinois House passed the AI Safety Measures Act 110-0, requiring annual third-party audits, incident reporting, and published safety frameworks for frontier AI developers above a compute threshold. The governor has signaled intent to sign, making it the first enacted US state law to mandate independent audits of AI labs.	With federal AI rules stalled, a single state has set a concrete compliance template that defines what documentation regulated buyers can now demand from AI vendors. The practical effect is a recognized yardstick for vendor due diligence and the beginning of a multi-state patchwork of differing requirements. IN PLAIN TERMS *Like the first state to require restaurants to post health-inspection grades in the window: one jurisdiction sets a standard, and soon every customer everywhere starts asking to see the grade before ordering.*	**A NEW VENDOR DUE DILIGENCE YARDSTICK FOR BANK AI PROCUREMENT** M&T's vendor risk and compliance teams now have a concrete external standard to cite when reviewing AI suppliers: independent audit results, incident-reporting history, and published safety frameworks. Any AI vendor serving the bank that operates above the compute threshold should be asked to demonstrate compliance or equivalent controls. Watch for additional state bills using Illinois as a template, which will require M&T's third-party risk program to track a patchwork rather than a single federal standard.
2 **FTC Fines Three Firms $930,000 for Falsely Marketing an AI Service That Did Not Exist** *Federal Trade Commission, May 21, 2026* The FTC settled with Cox Media Group, MindSift, and 1010 Digital Works for falsely marketing an AI-powered advertising service that used no AI at all, reselling data-broker email lists instead. Cox Media will pay $880,000, with future violations carrying penalties above $53,000 each.	Regulators are treating false AI capability claims as straightforward deceptive marketing, and enforcement does not require the technology to have worked. The action reinforces a parallel SEC trend pursuing AI-washing cases against financial firms. IN PLAIN TERMS *Like selling a home security system with cameras that were never plugged in: the regulator does not care that the footage was fake, only that you charged for protection you never provided.*	**AI-WASHING EXPOSURE REACHES BOTH VENDORS AND BANK MARKETING CLAIMS** M&T faces this risk on two fronts: any AI vendor whose capability claims cannot be documented creates procurement liability, and any customer-facing language describing bank services as AI-powered must match what those products actually do. Compliance and marketing should audit outward-facing AI claims against documented capability before the SEC applies the same lens it has used on investment advisers. The $53,000-per-violation structure means repeat exposure accumulates quickly.
3 **CNN Sues Perplexity Over Alleged Copyright Infringement of 17,000 News Stories** *CNN Business, Variety, NPR, Al Jazeera, May 28, 2026* CNN filed a federal copyright and trademark lawsuit against Perplexity AI, alleging unauthorized scraping of over 17,000 stories to power its products. Perplexity faces parallel suits from the New York Times, News Corp, and others, while OpenAI, Meta, and Google have struck licensing deals.	The AI content industry is splitting into licensed and litigation camps, and the outcome will determine what data can legally feed AI products and at what cost. For organizations using AI tools that surface third-party content, such as regulatory monitoring or market intelligence, vendor data provenance is now a procurement question. IN PLAIN TERMS *Like the early Napster lawsuits that forced the music industry to either license or sue: the eventual answer was both, and it reshaped how creators were compensated for the next 20 years.*	**AI VENDOR DATA PROVENANCE IS NOW A LOAN AND COMPLIANCE RESEARCH RISK** Any AI tool M&T uses for regulatory monitoring, market intelligence, or news summarization may be feeding on content obtained without license. Before renewing contracts with AI search or research vendors, procurement should require vendors to demonstrate licensed data provenance or confirm they are not among the parties facing active litigation. The bank's own legal exposure is indirect but real if a vendor's AI output is used in client-facing or compliance work built on disputed content.
4 **Anthropic Moves Claude Mythos Toward Release After Controlled Cyber Vetting** *Anthropic, Project Glasswing, cybersecurity press, May 26, 2026* Anthropic's Claude Mythos was withheld from release after it demonstrated the ability to find and chain zero-day exploits in real codebases. Applied defensively through Project Glasswing, it identified thousands of critical software flaws and is now moving toward a controlled commercial release for enterprise security use.	Frontier AI now has offensive-grade cybersecurity capability that is moving toward commercial availability for defenders, which also means comparable capability will reach attackers. Security planning should assume this class of tool is arriving on both sides of the threat surface. IN PLAIN TERMS *Like a powerful new tool that a manufacturer first keeps in-house, tests carefully, then sells only to licensed professionals: the careful path to market is itself a signal of how capable, and how double-edged, the tool is.*	**OFFENSIVE-GRADE AI VULNERABILITY TOOLS ARE APPROACHING THE MARKET: BANK SECURITY PLANS NEED TO CATCH UP** M&T's security and technology risk teams should treat the Mythos arc as a forward indicator: the same AI capability that found thousands of software vulnerabilities in a controlled setting will be available to threat actors within a similar timeframe. Existing penetration-testing assumptions and vendor vulnerability-disclosure programs should be reviewed against a threat model that includes AI-assisted zero-day chaining. This is a near-term planning item for the next security roadmap cycle, not a future horizon.
5 **Voice-Driven Workplaces Move From Experiment to Operating Model at AI-Native Firms** *The Information AI Agenda, May 27, 2026* At AI-first companies, staff are shifting from typing to continuous voice dictation using tools like Wispr, with offices adding sound-damping pods to handle acoustic and privacy load. The pattern is spreading and represents an early operating-model signal from AI-native firms.	If voice becomes the primary input for knowledge work, information security and HR policies built around keyboards and screens must extend to ambient speech capture, third-party transcription vendors, and which AI models hear which words. The productivity upside is real: dictation can roughly double knowledge-worker output. IN PLAIN TERMS *Like the trading floor going from open outcry to silent*	**AMBIENT VOICE CAPTURE CREATES NEW INFO-SEC AND COMPLIANCE OBLIGATIONS FOR BANK BRANCHES AND OFFICES** As M&T's own staff begin adopting AI voice tools, client conversations, loan terms, and compliance discussions previously conducted quietly at a keyboard may begin moving into ambient speech. The bank's information security policy, vendor agreements with transcription providers, and call-recording compliance obligations all need to be evaluated for widespread.

THE WEEK IN ONE LINE

This week's signals converge on a single theme: the regulatory and security perimeter for AI is hardening in concrete ways, not abstract ones. For M&T, the practical work is updating vendor due diligence criteria, auditing AI-washing exposure, and moving the bank's security planning horizon forward to account for offensive-grade AI tools arriving on both sides of the threat surface.